

[Open in app](#)

★ Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#) X

When One LLM Isn't Enough: How MA-RAG Uses a Team of Agents to Fix RAG's Biggest Problems

5 min read · Just now



Rishikeshavlal Patel

[Listen](#)[Share](#)[More](#)

Based on “MA-RAG: Multi-Agent Retrieval-Augmented Generation via Collaborative Chain-of-Thought Reasoning”

Link – <https://arxiv.org/pdf/2505.20096.pdf>



If you've played with modern chatbots, you've probably seen both sides of them:

- Ask a simple fact question → instant, correct answer.
- Ask something slightly messy like "*Who coached the team that beat Liverpool in the European Cup final where Jupp Heynckes played, and where was that match?*" → the model often guesses, confuses years, or just hallucinates.

Traditional Retrieval-Augmented Generation (RAG) was invented to fix this: instead of relying only on whatever the model memorized during training, the system retrieves documents from an external knowledge base (like Wikipedia) and feeds them to the LLM to ground its answer in real evidence.


In practice though, real-world questions are:

- **Ambiguous** (“the final” – which year exactly?)
- **Multi-hop** (first find the year, then find the city, then maybe the stadium)
- Wrapped in **noisy retrieval** (lots of irrelevant or semi-relevant text in the retrieved chunks).

So a single LLM plus a flat list of documents still struggles.

The MA-RAG paper proposes a different angle: Don’t treat RAG as a single black box. Treat it as a **team of specialized agents** that plan, retrieve, filter, and answer collaboratively.

In this article, I’ll walk through what MA-RAG is, how this multi-agent setup works, why it beats many existing RAG systems, and what it means for the future of agentic AI.




A 30-Second Recap: What Is “Vanilla” RAG?

Classic RAG looks like this:


1. You ask a question
2. A retriever finds the top-k relevant documents from a corpus (e.g., Wikipedia)
3. The LLM gets a prompt containing your question + those documents
4. It generates an answer

This setup already helps reduce hallucinations, but it has three big pain points for complex questions:

1. **One-shot retrieval** — The system gets one chance to retrieve documents. If the original query is vague or incomplete, retrieval may miss what we truly need.
2. **No real planning** — Multi-hop questions (“first find A, then use it to look up B”) are handled in one big step. The LLM is expected to both *plan* and *reason* and *answer* in a single pass.
3. **Noisy, oversized context** — Retrieved documents are often long and partially irrelevant. Dumping everything into the context window can confuse the model and waste tokens.



MA-RAG doesn't just tweak the retriever. It **restructures the entire pipeline** into a modular, agent-based workflow that explicitly tackles ambiguity, planning, and noise.



Enter MA-RAG: Turning RAG Into a Multi-Agent Team

1. Planner Agent — *the strategist*

- Reads the original question
- Detects ambiguity and complexity
- Breaks the question into a sequence of simpler **sub-tasks** (a “plan”)

2. Step Definer Agent — *the query engineer*

- Takes one step from the plan at a time
- Turns that abstract step into a concrete, detailed **sub-query**, using the original question + previous answers as context.

3. Retrieval Tool + Extractor Agent — *the librarian & filter*

- Retrieval tool fetches top-k documents (e.g., via FAISS dense retrieval).
- **Extractor Agent** reads those and pulls out only **the sentences or spans that are actually relevant to this step**.
- It outputs a small set of **clean, step-specific notes** instead of raw, noisy passages.

4. QA Agent — *the final storyteller*

- Takes the step’s sub-query + extracted notes.
- Produces an answer for that step.
- Step answers get passed forward and eventually combined into the final answer.

A Concrete Example: Jupp Heynckes and the European Cup Final

Step	What the agent focuses on	Example
1	Figure out which final we're talking about Planner + Step Definer + Retriever + Extractor + QA	In which
2	Use that year to find the location of the final Retriever + Extractor + QA	Answer: "Where wa
3	Combine everything into a final answer	Answer: "It was t

Let's go back to a slightly confusing football question:

"Where was the only European Cup Final in which Jupp Heynckes played held?"

MA-RAG instead turns this into a **mini-conversation between agents**. In the paper's example (also featured in my slides), the Planner creates this two-step plan:

Step 1: "Which year did Jupp Heynckes play in the European Cup Final?"

Step 2: "Where was the European Cup Final held in that year?"

Then the system runs:

- **Step Definer** sharpens Step 1 into a detailed search query
- **Retrieval Tool** gets documents about Jupp Heynckes and specific finals
- **Extractor** filters out everything except the part saying *he played in the 1977 European Cup Final*
- **QA Agent** answers: **1977**

Now Step 2 uses that answer:

- New sub-query: "*Where was the 1977 European Cup Final held?*"
- Retrieval + extraction find that the 1977 final was held in **Rome**
- QA Agent answers: **Rome**

Putting it all together, MA-RAG returns the final answer:

1977, in Rome.

Why MA-RAG Matters for the Future of Agentic AI

Stepping back, MA-RAG is interesting not only as "yet another RAG tweak," but as a **template for agentic systems**:

- It shows that you can treat RAG as a reasoning pipeline, not just an add-on to an LLM.
- It demonstrates that small models + good orchestration can beat larger models that work alone.
- It provides interpretable intermediate artifacts (plans, sub-queries, extracted notes) that humans can inspect for debugging and trust.

Final Thoughts

For me, the most compelling part of MA-RAG is not any single benchmark number, but the **shift in mindset**:

- From “let’s cram more tokens into a bigger context window”
- To “let’s coordinate smaller, focused reasoning steps across specialized agents”

As LLMs move from chatbots to **autonomous agents** solving real tasks — coding, research, analysis, planning — this style of design feels increasingly inevitable.

In my accompanying presentation and video walkthrough, I dive deeper into the actual agent prompts, example reasoning traces, and what it would look like to plug MA-RAG into a real application. But at a high level, the takeaway is simple:

One LLM is smart. A team of LLM agents, working together with a plan, can be much smarter.

[AI](#)[AI Agent](#)[Agentic Ai](#)[Multi Agent Ai](#)[Retrieval Augmented Gen](#)[Edit profile](#)

Written by Rishikeshavla Patel

0 followers · 1 following